

LECTURE 9

SQL II

Relational Databases and Querying Multiple Tables

Data 100/Data 200, Fall 2021 @ UC Berkeley

Fernando Pérez and Alvin Wan

(content by Alvin Wan, Anthony D. Joseph, Allen Shen, Josh Hug, John DeNero, Joseph Gonzalez)

LECTURE 9

Motivation

Normalization Definition and Takeaways

Data 100/Data 200, Fall 2021 @ UC Berkeley

Fernando Pérez and Alvin Wan

(content by Alvin Wan, Anthony D. Joseph, Allen Shen, Josh Hug, John DeNero, Joseph Gonzalez)

What & Why

Takeaways

How to Try SQL

Scene

id INT, PK	biome TEXT, NOT NULL	city TEXT, NOT NULL	visitors INT, >=0
1	desert	Las Vegas	100
2	freshwater	Honolulu	1000
3	freshwater	Las Vegas	100

Scene

id INT, PK	biome TEXT, NOT NULL	city TEXT, NOT NULL	city radius INT, >= 0	visitors INT, >=0
1	desert	Las Vegas	15	100
2	freshwater	Honolulu	12	1000
3	freshwater	Las Vegas	15	100

Why **redundancy** is a **bad** idea.

More Storage

Slower **INSERT**

Possible incongruities

Scene

id INT, PK	biome TEXT, NOT NULL	city TEXT, NOT NULL	city radius INT, >= 0	visitors INT, >=0
1	desert	Las Vegas	15	100
2	freshwater	Honolulu	12	1000
3	freshwater	Las Vegas	15	100

Scene

id INT, PK	biome TEXT, NOT NULL	city TEXT, NOT NULL	city radius INT, >= 0	visitors INT, >=0
1	desert	Las Vegas	15	100
2	freshwater	Honolulu	12	1000
3	freshwater	Las Vegas	15	100

Scene

id INT, PK	biome TEXT, NOT NULL	city_id INT, FK	visitors INT, >=0
1	desert	1	100
2	freshwater	2	1000
3	freshwater	1	100

City

id INT, PK	name TEXT, NOT NULL	radius INT, >=0
1	Las Vegas	15
2	Honolulu	12

Scene

id INT, PK	biome TEXT, NOT NULL	city_id INT, FK	visitors INT, >=0
1	desert	1	100
2	freshwater	2	1000
3	freshwater	1	100

many

City

id INT, PK	name TEXT, NOT NULL	radius INT, >=0
1	Las Vegas	15
2	Honolulu	12

one

Normalization is a database design technique to remove redundancy and anomalies.

Why **normalization** is a **good** idea.

More Compact

Faster **INSERT**

Impossible to have incongruities

Minimize redesign need

What & Why
Takeaways
How to Try SQL

What you should focus on.

Not memorize SQL syntax

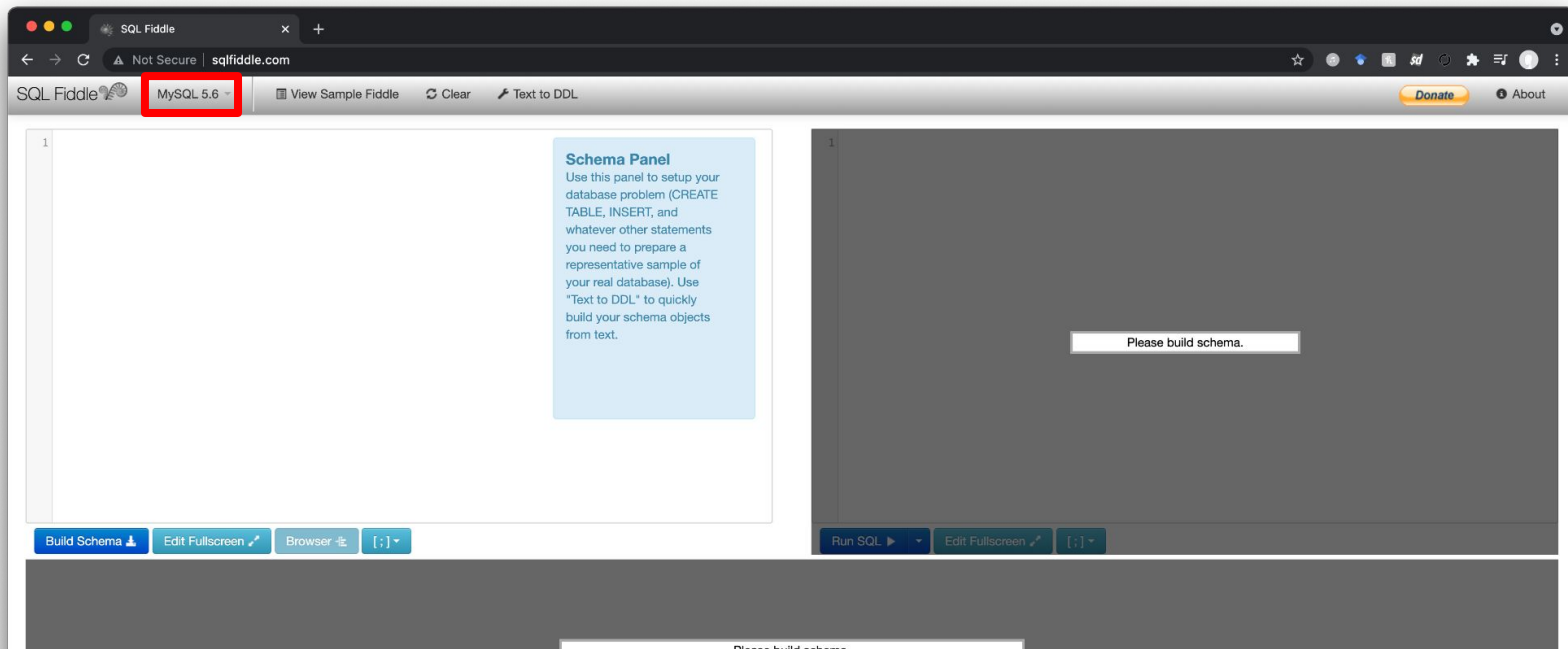
Understand what JOINS can do

Write clean SQL JOINS

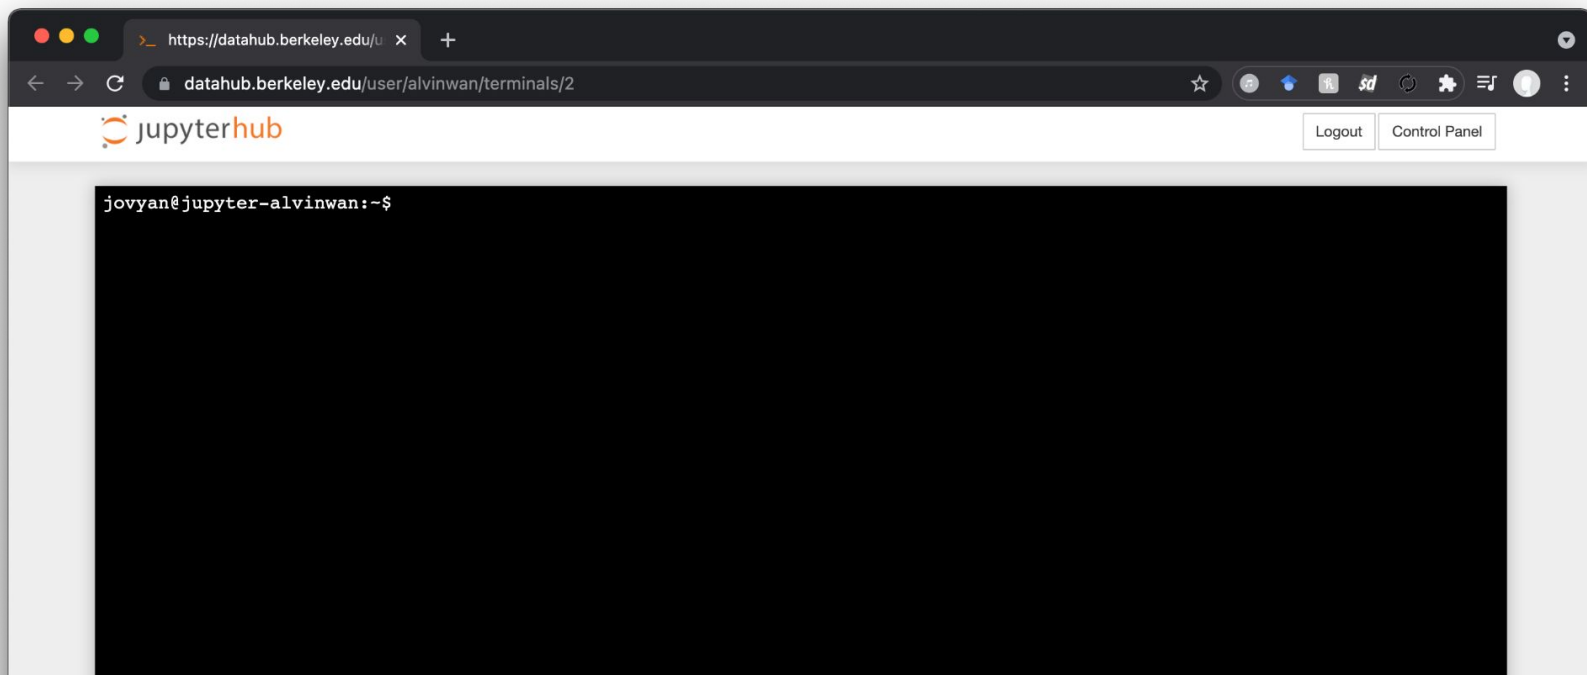
Debug SQL JOINS

What & Why Takeaways How to Try SQL

sqlfiddle.com



datahub.berkeley.edu terminal




Home Page - Select or create

datahub.berkeley.edu/user/alvinwan/tree?

☆

sd

jupyterhub

LogoutControl Panel

FilesRunningClustersNbextensions

Select items to perform actions on them.

Download DirectoryUploadNew

☐ 0

/

Name

Last Modified

File size

☐

SQL I Lecture.ipynb

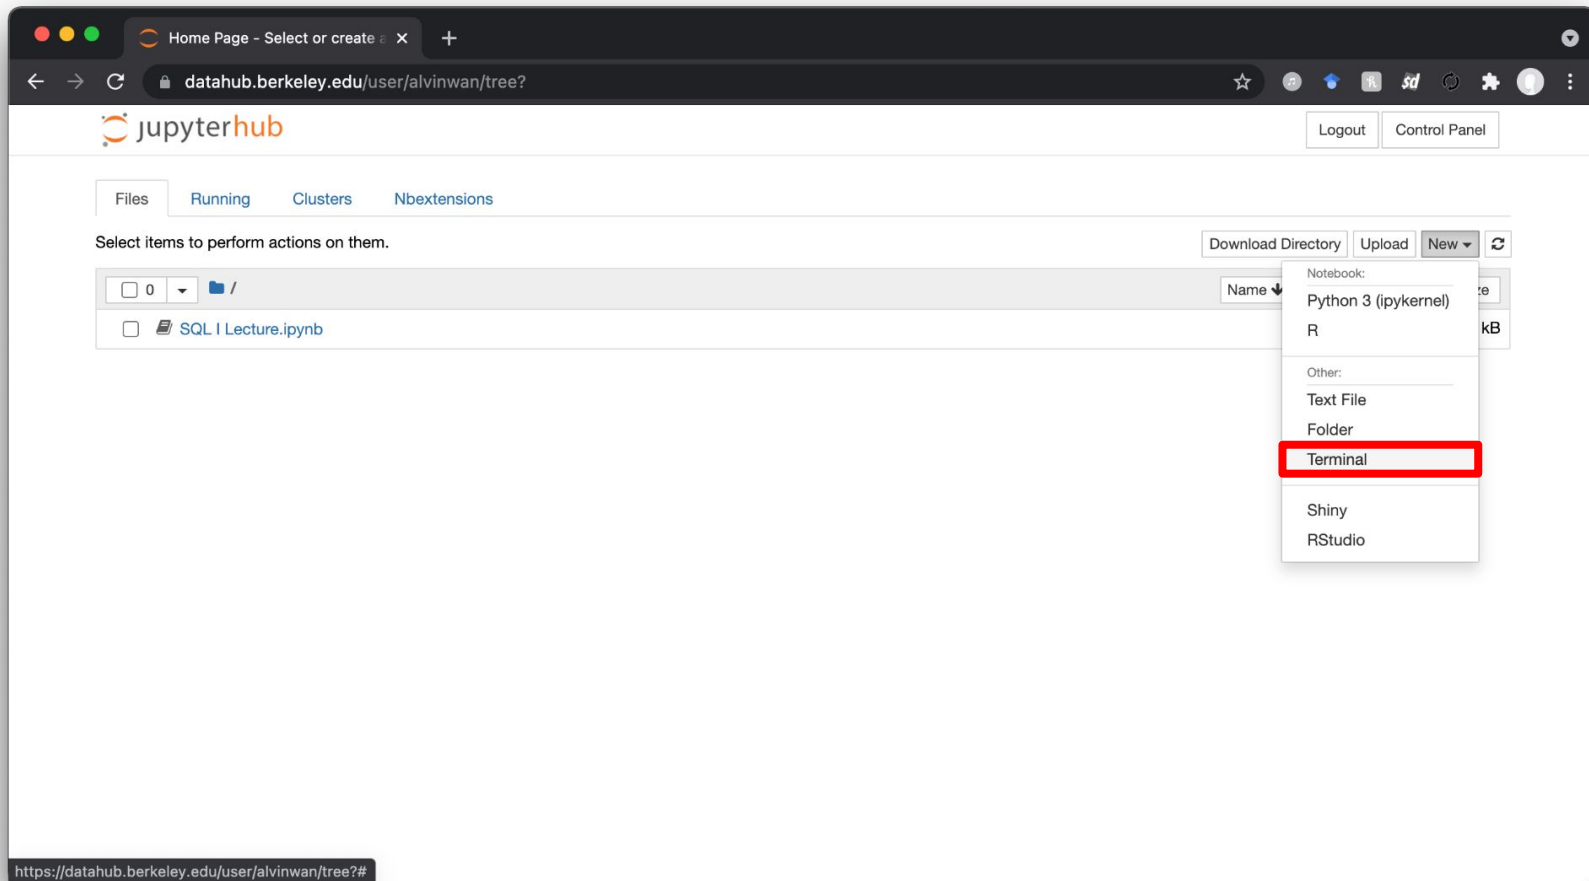
30 minutes ago


3.38 kB


BY

NC

SA







datahub.berkeley.edu notebook

SQL I Lecture - Jupyter Notebo x +

datahub.berkeley.edu/user/alvinwan/notebooks/SQL%20I%20Lecture.ipynb

jupyterhub SQL I Lecture Last Checkpoint: Last Friday at 5:06 PM (unsaved changes)

Logout Control Panel

File Edit View Insert Cell Kernel Widgets Help

Not Trusted Python 3 (ipykernel) O

Memory: 177.1 MB / 1 GB

```
In [4]: # Needed for now. Make sure to run this cell before running anything else in this notebook.
!pip install --quiet ipython-sql

Note: you may need to restart the kernel to use updated packages.
```

```
In [5]: !load_ext sql

The sql extension is already loaded. To reload it, use:
!reload_ext sql
```

```
In [7]: !!sql sqlite://
DROP TABLE IF EXISTS dragons;
CREATE TABLE dragons (name, birthyear);
INSERT INTO dragons VALUES ('drogon', 2011);
INSERT INTO dragons VALUES ('hiccup', 2019);
INSERT INTO dragons VALUES ('dragon 2', 2019);
```

Done.

SQL I Lecture - Jupyter Notebook

datahub.berkeley.edu/user/alvinwan/notebooks/SQL%20I%20Lecture.ipynb

jupyterhub SQL I Lecture Last Checkpoint: Last Friday at 5:06 PM (unsaved changes)

Logout Control Panel

File Edit View Insert Cell Kernel Widgets Help

Not Trusted Python 3 (ipykernel)

Memory: 225.8 MB / 1 GB

```
In [2]: %load_ext sql
```

```
In [ ]:
```

TAKEAWAY

Normalization **removes redundancy** and **anomalies** in your data. **Foreign keys** refer to **private keys** in a different table.

LECTURE 9

Types of Relationships

Many-to-One and Many-to-Many

Data 100/Data 200, Fall 2021 @ UC Berkeley

Fernando Pérez and Alvin Wan

(content by Alvin Wan, Anthony D. Joseph, Allen Shen, Josh Hug, John DeNero, Joseph Gonzalez)

Many-to-One
Many-to-Many

Scene

id INT, PK	biome TEXT, NOT NULL	city_id INT, FK	visitors INT, >=0
1	desert	1	100
2	freshwater	2	1000
3	freshwater	1	100

Many-to-One

City

id INT, PK	name TEXT, NOT NULL	radius INT, >=0
1	Las Vegas	15
2	Honolulu	12

Many Scenes to One City

Many-to-One
Many-to-Many

Scene

id INT, PK	biome TEXT, NOT NULL	city_id INT, FK	visitors INT, >=0
1	desert	1	100
2	freshwater	2	1000
3	freshwater	1	100

City

id INT, PK	name TEXT, NOT NULL	radius INT, >=0
1	Las Vegas	15
2	Honolulu	12

Scene

id INT, PK	biome TEXT, NOT NULL	city_id INT, FK	visitors INT, >=0
1	desert	1	100
2	freshwater	2	1000
3	freshwater	1	100

City

id INT, PK	name TEXT, NOT NULL	radius INT, >=0
1	Las Vegas	15
2	Honolulu	12

Biome

id INT, PK	name TEXT, NOT NULL
1	desert
2	freshwater

Scene

id INT, PK	biome_id INT, FK	city_id INT, FK	visitors INT, >=0
1	1	1	100
2	2	2	1000
3	2	1	100

City

id INT, PK	name TEXT, NOT NULL	radius INT, >=0
1	Las Vegas	15
2	Honolulu	12

One-to-Many

Biome

id INT, PK	name TEXT, NOT NULL
1	desert
2	freshwater

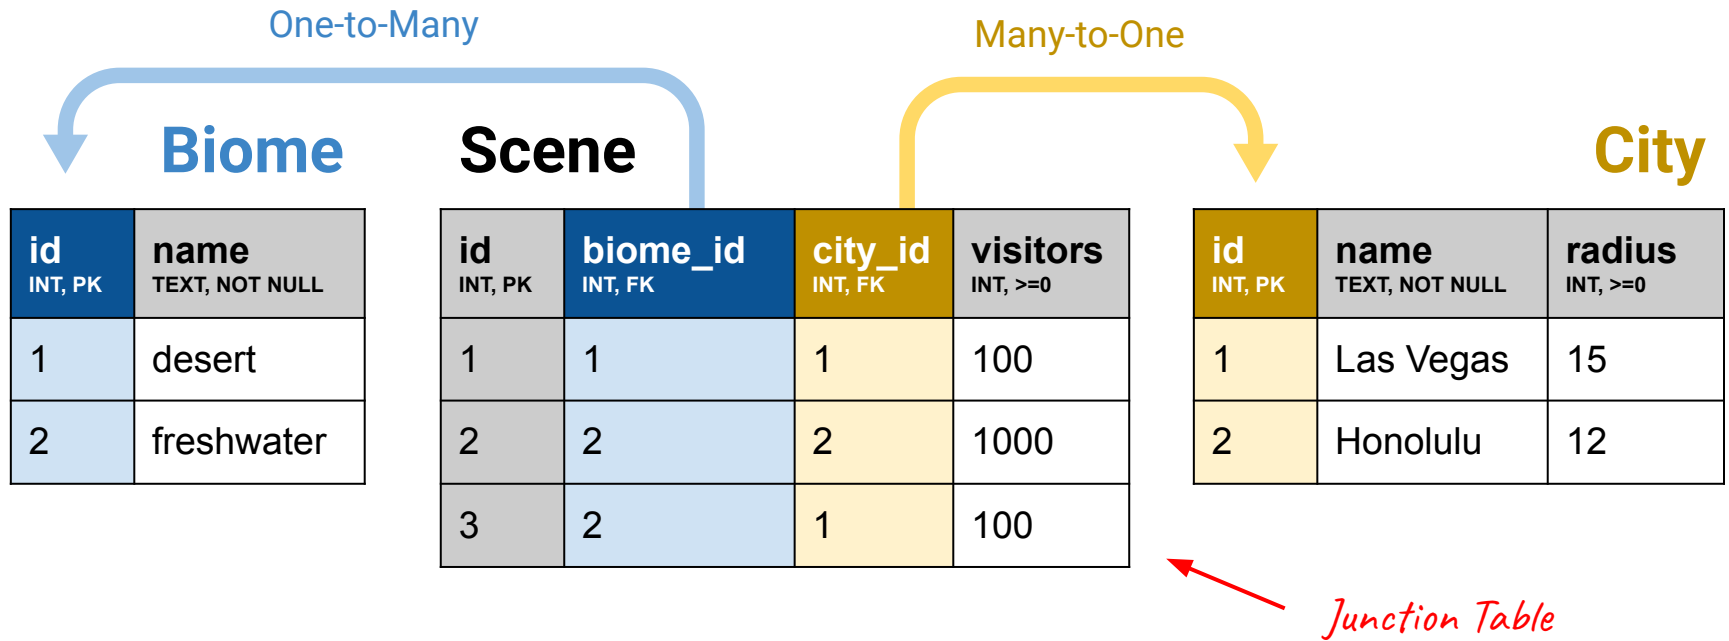
Scene

id INT, PK	biome_id INT, FK	city_id INT, FK	visitors INT, >=0
1	1	1	100
2	2	2	1000
3	2	1	100

City

id INT, PK	name TEXT, NOT NULL	radius INT, >=0
1	Las Vegas	15
2	Honolulu	12

One Biome to Many Scenes



Many Biomes to Many Cities

TAKEAWAY

Table relationships can be **one-to-many** or **many-to-many**. Use a junction table for many-to-many.

LECTURE 9

Conclusion

Takeaways

Data 100/Data 200, Fall 2021 @ UC Berkeley

Fernando Pérez and Alvin Wan

(content by Alvin Wan, Anthony D. Joseph, Allen Shen, Josh Hug, John DeNero, Joseph Gonzalez)

Summary

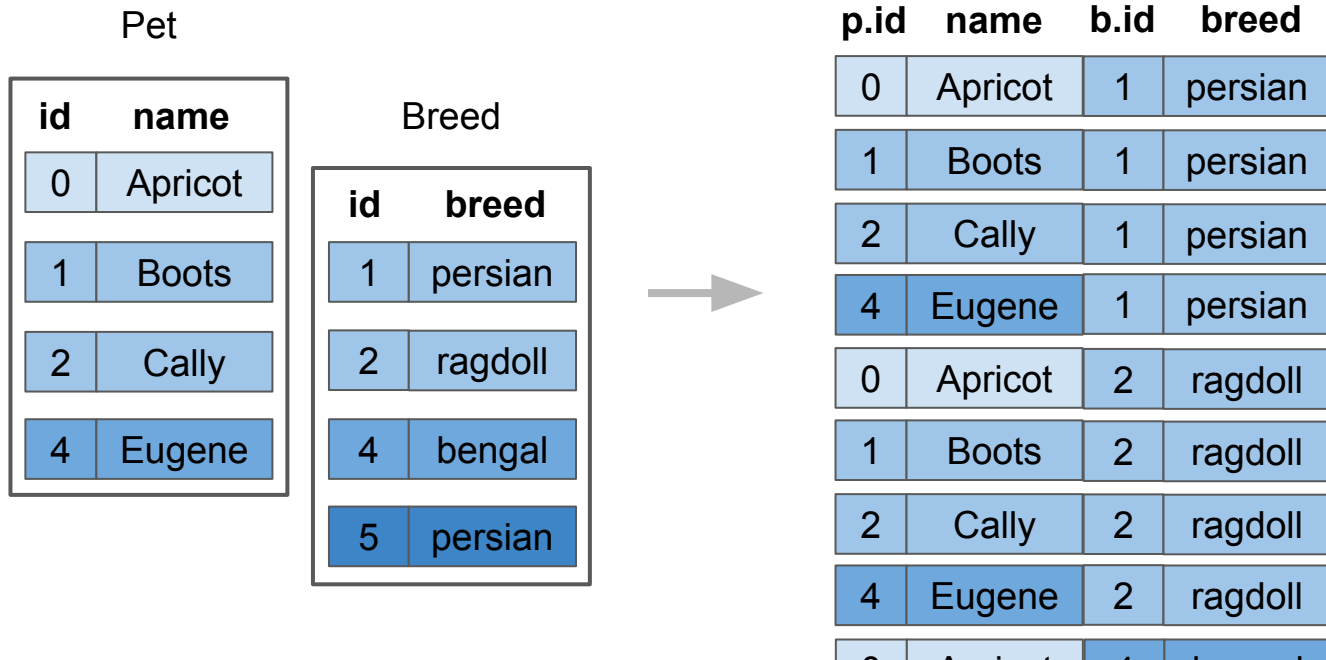
- Possible to subsample directly in SQL
- Common Table Expressions allow complex queries
 - Complexity not always needed
 - CTEs preferred over subqueries, for readability
- Use IS for NULL predicates
- Simpler queries are often possible!

Scene

id INT, PK	biome TEXT, NOT NULL	city TEXT, NOT NULL	visitors INT, >=0
0	desert	Las Vegas	100
1	marine	Honolulu	1000
2	freshwater	Paris	50
3	marine	Taipei	100
4	desert	Austin	25
5	freshwater	Austin	240
6	freshwater	Las Vegas	100

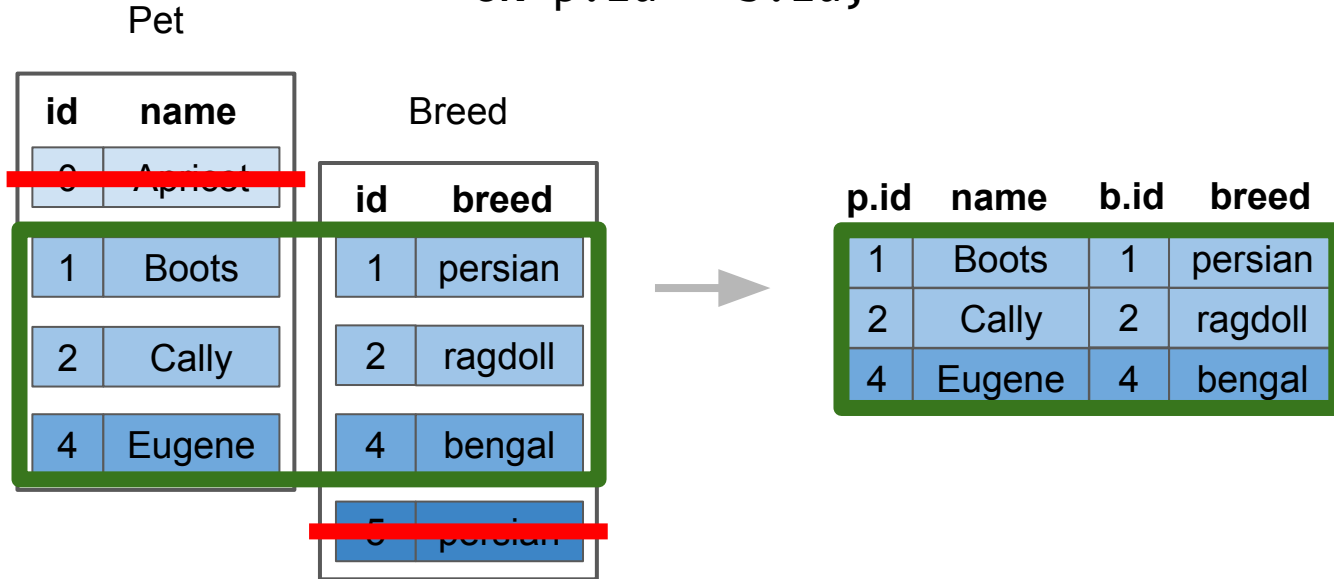
Cross Join: All Pairs of Rows

```
SELECT * FROM Pet AS p, Breed AS b;
```



Inner Join: Only Matching Rows

```
SELECT * FROM Pet AS p
JOIN Breed AS b
ON p.id = b.id;
```



Inner Join is Cross Join + Filter (conceptually)

```
SELECT *  
FROM Pet AS p, Breed AS b  
WHERE p.id = b.id;
```

Pet

id	name
0	Apricot
1	Boots
2	Cally
4	Eugene

Breed

id	breed
1	persian
2	ragdoll
4	bengal
5	persian

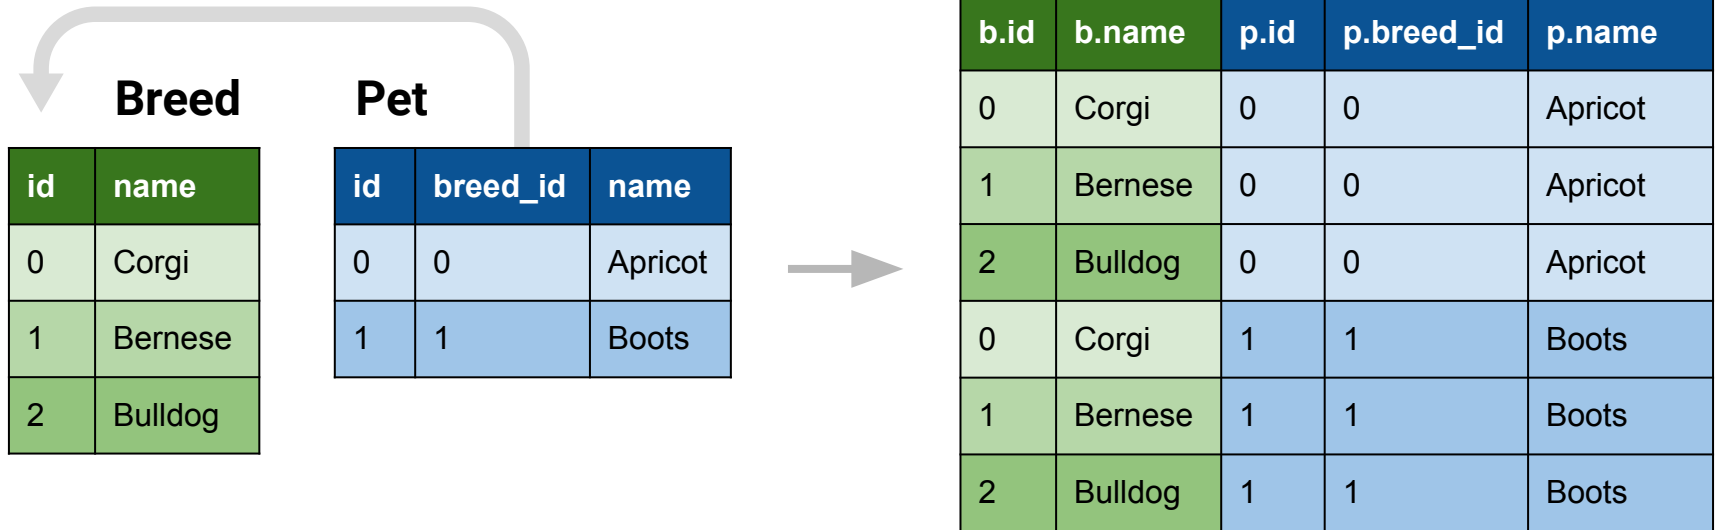


p.id	name	b.id	breed
0	Apricot	1	persian
1	Boots	1	persian
2	Cally	1	persian
4	Eugene	1	persian
0	Apricot	2	ragdoll
1	Boots	2	ragdoll
2	Cally	2	ragdoll
4	Eugene	2	ragdoll

Many-to-One
Many-to-Many
Cross Join
Inner Join
Outer Join
Join Conditions

Cross Join: All Pairs of Rows

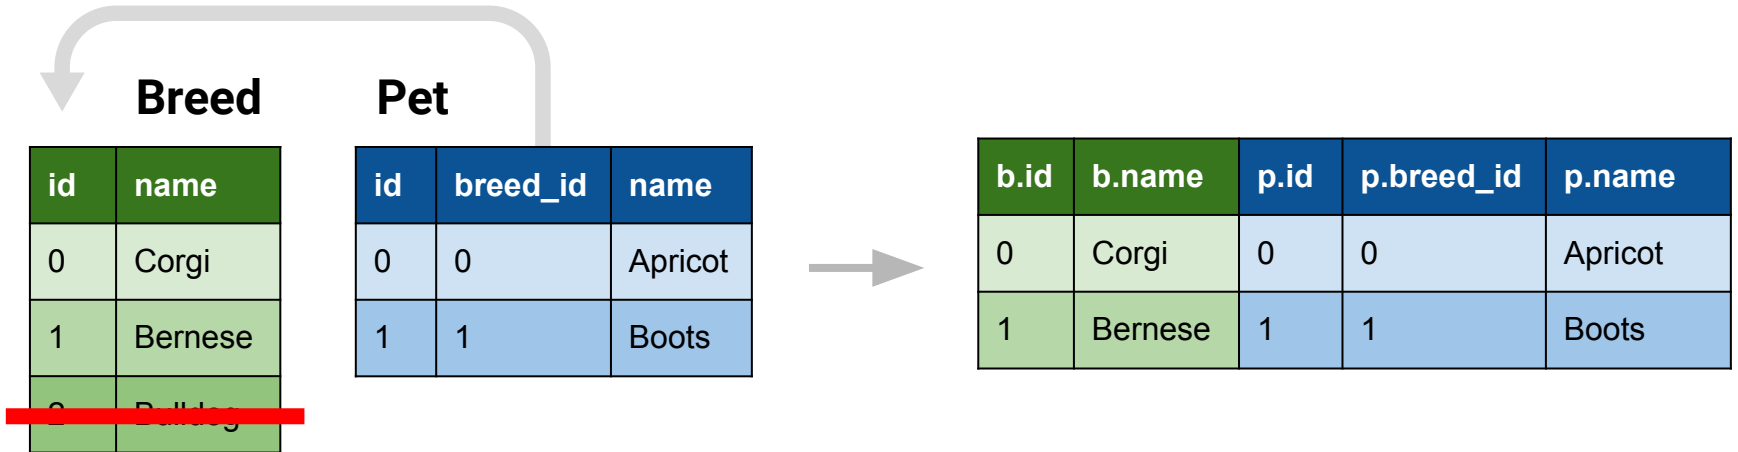
```
SELECT * FROM Pet AS p, Breed AS b;
```



Many-to-One
Many-to-Many
Cross Join
Inner Join
Outer Join
Join Conditions


Inner Join: Only Matching Rows

```
SELECT * FROM Pet AS p
JOIN Breed AS b
ON p.breed_id = b.id;
```



Inner Join is Cross Join + Filter (conceptually)

```
SELECT *  
FROM Pet AS p, Breed AS b  
WHERE p.breed_id = b.id;
```



id	name
0	Corgi
1	Bernese
2	Bulldog

id	breed_id	name
0	0	Apricot
1	1	Boots

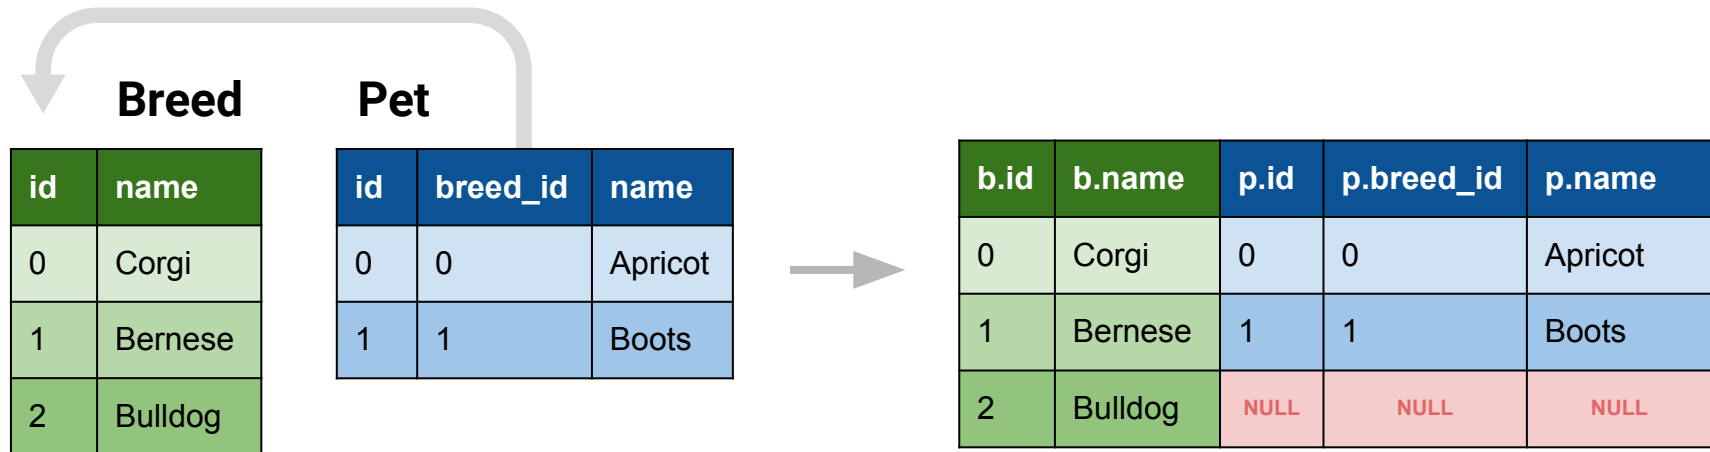


b.id	b.name	p.id	p.breed_id	p.name
0	Corgi	0	0	Apricot
1	Bernese	0	0	Apricot
2	Bulldog	0	0	Apricot
0	Corgi	1	1	Boots
1	Bernese	1	1	Boots
2	Bulldog	1	1	Boots

Many-to-One
Many-to-Many
Cross Join
Inner Join
Outer Join
Join Conditions

Left Outer Join: All Rows in 1st Table

```
SELECT * FROM Pet AS p
LEFT JOIN Breed AS b
ON p.id = b.id;
```



PRACTICAL TIP

Use **OUTER JOIN** to account for rows with no relationships.

Example: Report number of orders per user.
INNER JOIN would omit users with 0 orders.

Right Outer Join: All Rows in 2nd Table

NOT SUPPORTED
IN SQLITE

```
SELECT * FROM Pet AS p
RIGHT JOIN Breed AS b
ON p.id = b.id;
```

Pet

id	name
0	Apriest
1	Boots
2	Cally
4	Eugene

Breed

id	breed
1	persian
2	ragdoll
4	bengal
5	persian



p.id	name	b.id	breed
1	Boots	1	persian
2	Cally	2	ragdoll
4	Eugene	4	bengal
		5	persian

Missing values
are **NULL**.

Full Outer Join: All Rows in Both Tables

NOT SUPPORTED
IN SQLITE

```
SELECT * FROM Pet AS p
FULL OUTER JOIN Breed AS b
ON p.id = b.id;
```

Pet

id	name
0	Apricot
1	Boots
2	Cally
4	Eugene

Breed

id	breed
1	persian
2	ragdoll
4	bengal
5	persian



p.id	name	b.id	breed
0	Apricot		
1	Boots	1	persian
2	Cally	2	ragdoll
4	Eugene	4	bengal
		5	persian

Many-to-One
Many-to-Many
Cross Join
Inner Join
Outer Join
Join Conditions

Join conditions don't have to be equality

```
SELECT * FROM Student AS s  
JOIN Teacher AS t  
ON s.age > t.age;
```

Student

age	name
29	Jameel
37	Jian
20	Emma

Teacher

age	name
52	Ira
27	John
36	Anuja



s.age	s.name	t.age	t.name
29	Jameel	27	John
37	Jian	27	John
37	Jian	36	Anuja

TAKEAWAY

Use a junction table for many-to-many relationships. Use **OUTER JOINS** to account for rows with no relationships.